


Pendulum DQN: Clipping Hurts Performance (Seed 42, 20-episode moving average)



Q-value Violations Persist Throughout Training (50-62% of Q-values exceed $Q_{\max}=0$)

